

# EfficientNet Architecture Overview

give me the entire architecture of efficientnet

EfficientNet is a family of convolutional neural network architectures introduced by **Tan and Le (2019)** in the paper "[EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks](#)". It scales **depth**, **width**, and **resolution** using a **compound coefficient** to achieve better accuracy and efficiency.

---

## EfficientNet: Key Concepts

### 1. Compound Scaling

Instead of scaling dimensions arbitrarily, EfficientNet uses a **compound coefficient  $\phi$**  to scale:

- **Depth (d)** → more layers
- **Width (w)** → more channels
- **Resolution (r)** → larger input images

The scaling is defined as:

$$\begin{aligned}\text{depth: } d &= \alpha^\phi \\ \text{width: } w &= \beta^\phi \\ \text{resolution: } r &= \gamma^\phi \\ \text{subject to: } \alpha \cdot \beta^2 \cdot \gamma^2 &\approx 2\end{aligned}$$

---

## EfficientNet Architecture Overview

EfficientNet builds upon **MBConv blocks** (Mobile Inverted Bottleneck Convs) with:

- Depthwise separable convolutions
- Squeeze-and-Excitation blocks

- Swish activation (later SiLU)

The base model is **EfficientNet-B0**, and others like B1 to B7 are scaled versions.

## EfficientNet-B0 Architecture (Baseline)

Here is a detailed block-wise architecture for **EfficientNet-B0**:

Stage	Operator	Input Resolution	Channels (out)	Repeats	Kernel	Stride	SE Ratio
1	Conv3x3	224×224	32	1	3×3	2	-
2	MBConv1, k3x3	112×112	16	1	3×3	1	0.25
3	MBConv6, k3x3	112×112	24	2	3×3	2	0.25
4	MBConv6, k5x5	56×56	40	2	5×5	2	0.25
5	MBConv6, k3x3	28×28	80	3	3×3	2	0.25
6	MBConv6, k5x5	14×14	112	3	5×5	1	0.25
7	MBConv6, k5x5	14×14	192	4	5×5	2	0.25
8	MBConv6, k3x3	7×7	320	1	3×3	1	0.25
9	Conv1x1 + Pool + FC	7×7	1280	-	1×1	1	-

 **MBConv** stands for Mobile Inverted Bottleneck Convolution.

## MBConv Block Breakdown

Each MBConv block consists of:

1. Expansion (Conv 1×1)

2. Depthwise Convolution ( $k \times k$ )
3. SE block (Squeeze-and-Excitation)
4. Projection (Conv  $1 \times 1$ )
5. Residual connection (if input == output and stride = 1)

---

## EfficientNet Variants (B1–B7)

They are scaled versions of B0 using compound scaling:

Model	$\phi$	Resolution	Depth	Width	Parameters	Top-1 Accuracy
EfficientNet-B0	0	224×224	1.0×	1.0×	5.3M	77.1%
EfficientNet-B1	1	240×240	1.1×	1.0×	7.8M	79.1%
EfficientNet-B2	2	260×260	1.2×	1.1×	9.2M	80.1%
EfficientNet-B3	3	300×300	1.4×	1.2×	12M	81.6%
EfficientNet-B4	4	380×380	1.8×	1.4×	19M	83.0%
EfficientNet-B5	5	456×456	2.2×	1.6×	30M	83.6%
EfficientNet-B6	6	528×528	2.6×	1.8×	43M	84.0%
EfficientNet-B7	7	600×600	3.1×	2.0×	66M	84.4%

---

## Official and Reference Implementations

- Official TensorFlow implementation:

<https://github.com/tensorflow/tpu/tree/master/models/official/efficientnet>

(Highly recommended if you want paper-level fidelity)

- **PyTorch implementation** (popular):  
<https://github.com/lukemelas/EfficientNet-PyTorch>
  - **Keras (TF 2.x)**:  
<https://github.com/qubvel/efficientnet>
- 

## **Bonus: EfficientNet-Lite**

- Optimized for **mobile/edge devices**
  - Removed **swish activation** (replaced by ReLU6)
  - Smaller SE blocks
- 

If you'd like a **modular PyTorch or Keras code walkthrough of EfficientNet-B0**, I can provide that next. Would you like it layer-by-layer?

explain MBConv6, k3x3 w.r.t se ratio 0.25 completely

Absolutely. Let's **decompose and deeply understand** the **MBConv6, k3x3 block** with **Squeeze-and-Excitation (SE) ratio 0.25**, which is a **core building block of EfficientNet**.

---

## **What is MBConv?**

**MBConv** stands for **Mobile Inverted Bottleneck Convolution**, first introduced in MobileNetV2 and later enhanced with **SE blocks** and **Swish activation** in EfficientNet.

---

## **MBConv6, k3x3 (with SE ratio 0.25)**

### **Specifications:**

- **6** → Expansion factor (input channels × 6)
  - **k3x3** → Kernel size of the depthwise convolution
  - **SE ratio = 0.25** → SE bottleneck channels = 0.25 × expanded channels
- 

## Full Architecture of MBConv6, k3x3

Assume:

- Input tensor: shape (**B, H, W, C**) — i.e., batch × height × width × input channels
- Let's say **C = 32** (input channels)

Now step-by-step:

---

### 1. Expansion Phase (Pointwise Conv1x1)

- Expand channels:

$$\text{Expanded Channels} = C_{exp} = C \times 6 = 32 \times 6 = 192$$

- Apply **1x1 convolution** to map from 32 → 192 channels
  - Shape: (**B, H, W, 192**)
  - Followed by:
    - **BatchNorm**
    - **Swish (SiLU)** activation
- 

### 2. Depthwise Convolution (Kernel = 3×3)

- Apply a **depthwise separable convolution** (1 filter per input channel, not across channels)
- Kernel size = 3×3, **stride = 1 or 2**

- Output shape: still **(B, H, W, 192)** (if stride = 1)
  - Followed by:
    - **BatchNorm**
    - **Swish activation**
- 

### 3. 💡 **Squeeze-and-Excitation (SE) Block**

This improves channel-wise attention. Here's the breakdown:

#### a. 📦 **Squeeze**

- Perform **Global Average Pooling** → output shape: **(B, 192, 1, 1)** → flattened to **(B, 192)**  
(One scalar per channel)

#### b. 🔄 **Excitation**

- First Dense Layer:

$$\text{Reduced Channels} = 192 \times 0.25 = 48$$

- Linear (FC) → shape becomes (B, 48)
- Activation: **Swish**
- Second Dense Layer:
  - Linear → back to (B, 192)
  - Activation: **Sigmoid**
- Output: channel-wise scale factors in range (0, 1)

#### c. 🔄 **Recalibrate**

- Multiply these 192 scale factors back with the depthwise output (broadcasted across H and W):

$$\text{Output shape still: } (B, H, W, 192)$$

---

### 4. 📦 **Projection Phase (Pointwise Conv1x1)**

- **1x1 convolution** to project back to original channel size: **192 → 32**
- Shape: **(B, H, W, 32)**
- Followed by:
  - **BatchNorm**
  - **No activation**

## 5. Residual Connection (Skip Connection)

- If **stride = 1** and **input\_channels == output\_channels**:

$$\text{Output} = \text{Input} + \text{Projected Output}$$

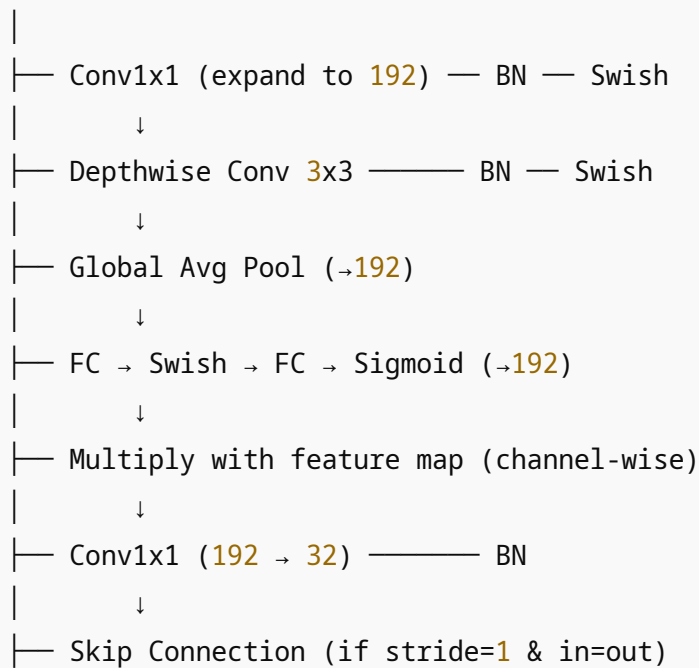
- Else → no skip connection

## Summary Table for MBConv6, k3x3, SE=0.25

Phase	Operation	Channels (in→out)	Notes
Expansion	1×1 Conv + BN + Swish	32 → 192	Expand channels
Depthwise	3×3 DW Conv + BN + Swish	192 → 192	Kernel = 3×3
SE Block	GAP → FC(48) → FC(192)	192 → 48 → 192	SE ratio = 0.25
Projection	1×1 Conv + BN	192 → 32	Reduce back to original dim
Residual	Add input (if possible)	32 + 32 → 32	Only if stride = 1

## Visual Diagram

Input (B, H, W, 32)



Output (B, H, W, 32)

## PyTorch Code Reference (Simplified)

python

```
class MBConv(nn.Module):
    def __init__(self, in_channels, out_channels, expand_ratio=6, se_ratio=0.25,
kernel_size=3, stride=1):
        super().__init__()
        mid_channels = in_channels * expand_ratio

        self.use_residual = (in_channels == out_channels and stride == 1)

        self.expand = nn.Sequential(
            nn.Conv2d(in_channels, mid_channels, 1, bias=False),
            nn.BatchNorm2d(mid_channels),
            nn.SiLU()
        ) if expand_ratio != 1 else nn.Identity()

        self.dwconv = nn.Sequential(
```



```

        nn.Conv2d(mid_channels, mid_channels, kernel_size, stride,
                  padding=kernel_size//2, groups=mid_channels, bias=False),
        nn.BatchNorm2d(mid_channels),
        nn.SiLU()
    )

    se_channels = int(mid_channels * se_ratio)
    self.se = nn.Sequential(
        nn.AdaptiveAvgPool2d(1),
        nn.Conv2d(mid_channels, se_channels, 1),
        nn.SiLU(),
        nn.Conv2d(se_channels, mid_channels, 1),
        nn.Sigmoid()
    )

    self.project = nn.Sequential(
        nn.Conv2d(mid_channels, out_channels, 1, bias=False),
        nn.BatchNorm2d(out_channels)
    )

def forward(self, x):
    identity = x
    x = self.expand(x)
    x = self.dwconv(x)
    x = x * self.se(x)
    x = self.project(x)
    return x + identity if self.use_residual else x

```

Would you like this written in **TensorFlow/Keras** too, or want me to help you visualize **all MBConv blocks in EfficientNet-B0 layer-by-layer with shapes**?

yes

Great! Let's now go **layer-by-layer through EfficientNet-B0**, explaining each **MBConv block with shapes, SE application**, and **what happens at each step**, so you can visualize the architecture clearly and code it from scratch.



# EfficientNet-B0 Detailed Layer-by-Layer Architecture

Input: 224×224×3

SE Ratio: 0.25

Activation: Swish

Note: MBConv6 → expansion ratio = 6, MBConv1 → no expansion

## ✓ Stage 0: Stem

Layer	Type	Output Shape	Notes
Conv2D	3×3, stride=2	112×112×32	Initial conv layer
BN + Swish	-	112×112×32	

## ✓ Stage 1: MBConv1, k3x3 (No expansion)

Layer	Type	Output Shape	Notes
MBConv1	k=3, stride=1	112×112×16	No expansion, SE applied, 1 block

### 🔄 SE Block:

- Squeeze → 16 → FC → 4 → FC → 16 → Sigmoid
- Scale: (B, 112, 112, 16) × (B, 1, 1, 16)

## ✓ Stage 2: MBConv6, k3x3, 2 blocks

Layer	Type	Output Shape	Notes
MBConv6_1	stride=2	56×56×24	Expand 16→96→SE→96→24, no residual
MBConv6_2	stride=1	56×56×24	Residual connection enabled

#### 🔄 SE Block per block:

- Squeeze → 96 → FC → 24 → FC → 96 → Sigmoid

### ✅ Stage 3: MBConv6, k5x5, 2 blocks

Layer	Type	Output Shape	Notes
MBConv6_1	stride=2	28×28×40	Expand 24→144→SE→144→40
MBConv6_2	stride=1	28×28×40	Residual enabled

#### 🔄 SE Block:

- SE channels =  $144 \times 0.25 = 36$

### ✅ Stage 4: MBConv6, k3x3, 3 blocks

Layer	Type	Output Shape	Notes
MBConv6_1	stride=2	14×14×80	Expand 40→240→SE→240→80
MBConv6_2	stride=1	14×14×80	Residual
MBConv6_3	stride=1	14×14×80	Residual

#### 🔄 SE Block:

- SE channels =  $240 \times 0.25 = 60$

## ✓ Stage 5: MBConv6, k5x5, 3 blocks

Layer	Type	Output Shape	Notes
MBConv6_1	stride=1	14×14×112	Expand 80→480→SE→480→112
MBConv6_2	stride=1	14×14×112	Residual
MBConv6_3	stride=1	14×14×112	Residual

### SE Block:

- SE channels =  $480 \times 0.25 = 120$

## ✓ Stage 6: MBConv6, k5x5, 4 blocks

Layer	Type	Output Shape	Notes
MBConv6_1	stride=2	7×7×192	Expand 112→672→SE→672→192
MBConv6_2/3/4	stride=1	7×7×192	Residual

### SE Block:

- SE channels =  $672 \times 0.25 = 168$

## ✓ Stage 7: MBConv6, k3x3, 1 block

Layer	Type	Output Shape	Notes
MBConv6	stride=1	7×7×320	Expand 192→1152→SE→1152→320

### SE Block:

- SE channels =  $1152 \times 0.25 = 288$

✓ Stage 8: Head

Layer	Type	Output Shape	Notes
Conv2D	1×1, 320→1280	7×7×1280	Final conv layer before pooling
BN + Swish	-	7×7×1280	
Global AvgPool	-	1×1×1280	Reduces spatial dims
Dropout	-	1280	Dropout(p=0.2)
FC	-	1000	ImageNet classifier

🧱 Full Block Table Summary (EfficientNet-B0)

Stage	Block Type	Repeats	Input Res	Output Channels	Expansion	Kernel	Stride	SE Channels
0	Conv3×3	1	224×224	32	-	3×3	2	-
1	MBConv1	1	112×112	16	1	3×3	1	4
2	MBConv6	2	112×112	24	6	3×3	2,1	24
3	MBConv6	2	56×56	40	6	5×5	2,1	36
4	MBConv6	3	28×28	80	6	3×3	2,1	60
5	MBConv6	3	14×14	112	6	5×5	1	120
6	MBConv6	4	14×14	192	6	5×5	2,1	168

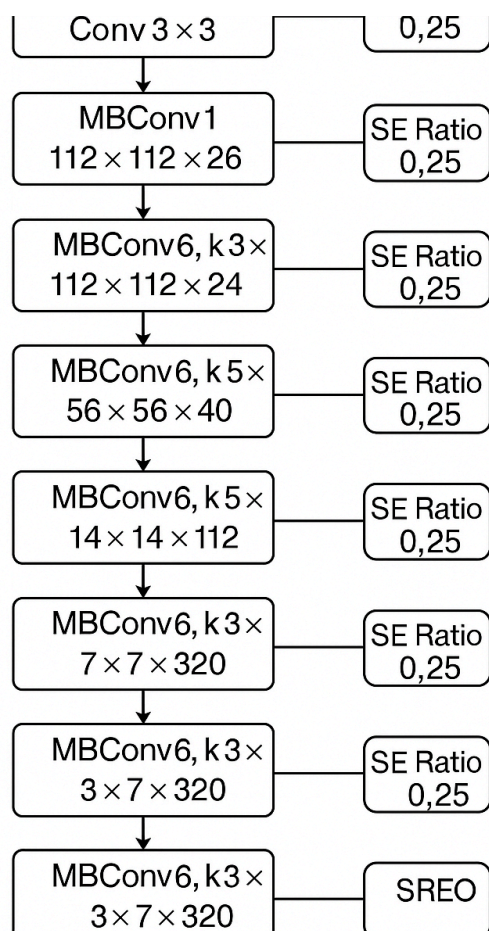
Stage	Block Type	Repeats	Input Res	Output Channels	Expansion	Kernel	Stride	SE Channels
7	MBConv6	1	7×7	320	6	3×3	1	288
8	Conv1x1 + Pool	1	7×7	1280	-	1×1	1	-

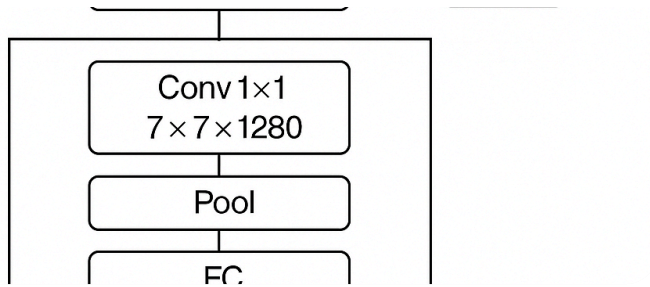
Would you like me to now:

- 🧱 Generate a **full Keras or PyTorch class for EfficientNet-B0** (modular)?
- 📊 Visualize the block diagram?
- 🖋️ Write minimal reproducible MBConv code with SE block?

Let me know which direction you want next!

Visualize the block diagram?





why even need squeeze excitation block

Great question. The **Squeeze-and-Excitation (SE) block** plays a critical role in **channel-wise attention**, which enhances the representational power of a CNN without significant computational cost. Let's break it down:

## Objective:

To allow the network to **learn which channels are more important** for a given task and **reweight them adaptively**.

## Intuition Behind SE Block

CNNs process features spatially and across channels, but **not all channels contribute equally** to the final prediction. Some may represent edges, others colors, textures, etc. SE blocks **learn to emphasize** the most informative channels while **suppressing less useful** ones.

Think of it as a **soft attention mechanism** over the channels.

## How It Works: (Step-by-Step in MBConv)

Assume feature map  $X$  of shape  $(B, H, W, C)$

### 1. Squeeze

- Global Average Pooling → shape becomes (B, C)
- Captures **global context** per channel (like summary stats)

$$z_c = \frac{1}{H \cdot W} \sum_{i=1}^H \sum_{j=1}^W X_{i,j,c}$$

## 2. Excitation

- Pass through 2 FC layers:
  - FC1 reduces dimensions →  $C \rightarrow C \times r$  (usually  $r=0.25$ )
  - Activation: Swish or ReLU
  - FC2 expands →  $C \times r \rightarrow C$
  - Activation: **Sigmoid** → values between 0 and 1 (soft gates)

This gives a vector  $\mathbf{s} \in \mathbb{R}^c$  of weights.

## 3. Scale

- Multiply original feature map channel-wise with  $\mathbf{s}$
- Broadcast  $\mathbf{s}$  across spatial dimensions

$$\tilde{X}_c = s_c \cdot X_c$$



## Why It's Needed (Benefits)

Benefit	Description
<b>Adaptive recalibration</b>	Learns which channels are important and adjusts dynamically per input
<b>Boosts performance</b>	Consistently improves accuracy on image classification, detection, etc.
<b>Minimal cost</b>	Adds very few parameters and FLOPs (esp. with bottleneck reduction ratio $r=0.25$ )
<b>Encodes global context</b>	Squeeze step captures context beyond local receptive field
<b>Plug-and-play</b>	Can be added to almost any CNN block (ResNet, MobileNet, etc.)



---


## Real-world Analogy

Imagine you're analyzing a news article with multiple sections: Politics, Tech, Sports, Health.

SE block is like a reader who first **skims the whole article** (squeeze), then decides, **"Tech and Health seem more relevant today,"** (excitation), and focuses more on those (scale).

---

## Evidence in Research

- Introduced in: **"Squeeze-and-Excitation Networks"** (Hu et al., 2018)  
 [Paper](#)
  - Achieved **ImageNet top-1 accuracy of 85.1%** on SENet-154
  - Now used in:
    - EfficientNet (SE + MBConv)
    - ResNeSt (Split-Attention)
    - MobileNetV3 (Lite SE)
    - YOLOv5 (optional)
- 

## Without SE block

If you **remove SE**, the network treats all channels equally. It still learns discriminative filters, but lacks an **explicit mechanism** to modulate their importance **based on global context** of the input.

---

## When is SE most effective?

- When channels represent **semantically diverse information**
  - In **deeper** models (EfficientNet, ResNeXt, etc.)
  - Tasks requiring **fine-grained discrimination** (e.g., medical imaging, Re-ID)
- 

Would you like me to show:

- Minimal PyTorch/Keras code for SE block?
- Visualization of how it adjusts channel activations before/after?

Let me know how deep you want to go into SE mechanics.

simple code in keras with comments

Here's a **simple and clean Keras implementation of a Squeeze-and-Excitation (SE) block**, with detailed comments explaining each step. You can plug this into any CNN block like MBConv, ResNet, etc.

---

## ✓ Keras Implementation of SE Block

python

```
from tensorflow.keras import layers

def se_block(input_tensor, se_ratio=0.25):
    """
    Squeeze-and-Excitation Block

    Args:
        input_tensor: Input feature map (B, H, W, C)
        se_ratio: Reduction ratio for bottleneck (default = 0.25)

    Returns:
        Tensor with SE channel-wise recalibration applied
    """
    channels = input_tensor.shape[-1] # Get the number of channels (C)
```

```

reduced_channels = max(1, int(channels * se_ratio)) # C * 0.25

# Squeeze: Global Average Pooling → (B, C)
se = layers.GlobalAveragePooling2D()(input_tensor)

# Reshape to (B, 1, 1, C) to match input spatial dims
se = layers.Reshape((1, 1, channels))(se)

# Excitation: FC1 (reduce) → Swish → FC2 (expand) → Sigmoid
se = layers.Conv2D(reduced_channels, kernel_size=1, activation='swish')(se)
se = layers.Conv2D(channels, kernel_size=1, activation='sigmoid')(se)

# Scale: Multiply original input by the excitation weights
output_tensor = layers.Multiply()(input_tensor, se)

return output_tensor

```

## Example Usage in a Conv Block:

python

```



def example_conv_block_with_se(input_tensor):
    x = layers.Conv2D(64, 3, padding='same', activation='swish')(input_tensor)
    x = layers.BatchNormalization()(x)


    # Apply SE block
    x = se_block(x, se_ratio=0.25)

    return x

```

Let me know if you want:

-  Integration inside **MobileNet/ResNet/EfficientNet**
-  Visualization of SE outputs

-  Variant: Lite SE (used in MobileNetV3)

Just say the word!

code MBConv6, k3x3, SE=0.25

Here's a complete and clean **Keras implementation of the** MBConv6, k3x3, SE=0.25 **block** — as used in EfficientNet. It includes:

- **Expansion (Conv1x1)**
- **Depthwise convolution (3x3)**
- **Squeeze-and-Excitation block (ratio=0.25)**
- **Projection (Conv1x1)**
- **Residual connection (if applicable)**

### ✓ mbconv\_block() in Keras (MBConv6, k3x3, SE=0.25)

python

```
from tensorflow.keras import layers

def se_block(input_tensor, se_ratio=0.25):
    """Squeeze-and-Excitation block"""
    channels = input_tensor.shape[-1]
    reduced_channels = max(1, int(channels * se_ratio))

    se = layers.GlobalAveragePooling2D()(input_tensor)
    se = layers.Reshape((1, 1, channels))(se)
    se = layers.Conv2D(reduced_channels, kernel_size=1, activation='swish')(se)
    se = layers.Conv2D(channels, kernel_size=1, activation='sigmoid')(se)
    return layers.Multiply()([input_tensor, se])
```

python

```

def mbconv_block(inputs, in_channels, out_channels, kernel_size=3, stride=1,
expand_ratio=6, se_ratio=0.25):
    """
    MBConv6 Block with:
    - Expansion
    - Depthwise Convolution (3x3)
    - Squeeze-and-Excitation
    - Projection
    - Optional Residual
    """
    # === Expansion phase ===
    x = inputs
    expanded_channels = in_channels * expand_ratio

    if expand_ratio != 1:
        x = layers.Conv2D(expanded_channels, kernel_size=1, padding='same',
use_bias=False)(x)
        x = layers.BatchNormalization()(x)
        x = layers.Activation('swish')(x)

    # === Depthwise convolution ===
    x = layers.DepthwiseConv2D(kernel_size=kernel_size, strides=stride,
padding='same', use_bias=False)(x)
    x = layers.BatchNormalization()(x)
    x = layers.Activation('swish')(x)

    # === Squeeze-and-Excitation ===
    x = se_block(x, se_ratio=se_ratio)

    # === Projection phase ===
    x = layers.Conv2D(out_channels, kernel_size=1, padding='same', use_bias=False)
(x)
    x = layers.BatchNormalization()(x)

    # === Residual connection ===
    if stride == 1 and in_channels == out_channels:
        x = layers.Add()([x, inputs])

    return x

```

## Example Usage:

python

```
from tensorflow.keras import Input, Model

# Input shape: (Batch, H, W, Channels)
inp = Input(shape=(56, 56, 24))
x = mbconv_block(inp, in_channels=24, out_channels=24, kernel_size=3, stride=1,
expand_ratio=6, se_ratio=0.25)
model = Model(inputs=inp, outputs=x)
model.summary()
```

---

## Notes:

- You can change `kernel_size=5` for `MBConv6`, `k5x5`.
- `stride=2` for downsampling.
- Use this inside a loop to build EfficientNet stages.
- For expansion ratio = 1, expansion phase is skipped.

---

Want the full **EfficientNet-B0** model built using these blocks next? Or a version with skip/residual visualization?