

ReLU Backpropagation Derivatives

explaining to me using derivatives relu backpropagation using example.

Understanding ReLU Backpropagation Using Derivatives

ReLU (Rectified Linear Unit) is a widely used activation function in neural networks, defined as:

$$\text{ReLU}(x) = \max(0, x)$$

The derivative of ReLU plays a crucial role in backpropagation, which is used to update weights in a neural network.

Step 1: Understanding ReLU and Its Derivative

The ReLU function is:

$$f(x) = \begin{cases} x, & \text{if } x > 0 \\ 0, & \text{if } x \leq 0 \end{cases}$$

Derivative of ReLU

The derivative of ReLU is:

$$f'(x) = \begin{cases} 1, & \text{if } x > 0 \\ 0, & \text{if } x \leq 0 \end{cases}$$

This means:

- If x is positive, the gradient is **1** (ReLU behaves like a linear function).
 - If x is negative or zero, the gradient is **0** (ReLU blocks the gradient, leading to the **dying ReLU problem** in some cases).
-

Step 2: Backpropagation with ReLU Example

Let's go through an example of a **simple neural network** with one hidden layer.

Network Structure

Consider a **single neuron** with:

- Input $x = 2$
- Weight $w = 3$
- Bias $b = 1$

- Activation function: ReLU

The forward pass computes:

$$z = wx + b = (3)(2) + 1 = 7$$

Applying ReLU:

$$a = \max(0, 7) = 7$$

Assume the loss function is:

$$L = (a - y)^2$$

where $y = 5$ (target output).

Step 3: Computing Gradients Using Backpropagation

To update w using **gradient descent**, we need:

$$\frac{dL}{dw}$$

Using the chain rule:

$$\frac{dL}{dw} = \frac{dL}{da} \cdot \frac{da}{dz} \cdot \frac{dz}{dw}$$

1. Compute $\frac{dL}{da}$

$$L = (a - y)^2$$

$$\frac{dL}{da} = 2(a - y) = 2(7 - 5) = 4$$

2. Compute $\frac{da}{dz}$ (Derivative of ReLU)

Since $z = 7$ (which is **positive**), we use:

$$\frac{da}{dz} = 1$$

3. Compute $\frac{dz}{dw}$

$$\frac{dz}{dw} = x = 2$$

Step 4: Update the Weight

$$\frac{dL}{dw} = 4 \times 1 \times 2 = 8$$

Using gradient descent with learning rate $\eta = 0.1$:

$$w_{\text{new}} = w - \eta \cdot \frac{dL}{dw}$$
$$w_{\text{new}} = 3 - 0.1 \times 8 = 2.2$$

Key Observations

1. **ReLU keeps gradients alive for positive values**, preventing vanishing gradients.
2. **For negative values, ReLU sets gradients to 0**, meaning neurons that get stuck in the negative region stop learning.
3. **Weight updates depend on the derivative of ReLU**, which is either **1 or 0**, leading to efficient computation.

Would you like a more complex example with a full backpropagation through multiple layers? 🚀

how would sigmoid do in this example instead of relu

Backpropagation with Sigmoid vs. ReLU: A Comparison Using the Same Example

Now, let's go through the same **forward and backward pass** using the **Sigmoid activation function** instead of ReLU and compare the results.

Step 1: Define the Sigmoid Activation Function

The Sigmoid function is:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

Its derivative is:

$$\sigma'(z) = \sigma(z)(1 - \sigma(z))$$

Unlike ReLU, Sigmoid **never gives a zero gradient**, but it **saturates** when z is too large or too small, leading to **vanishing gradients**.

Step 2: Forward Pass with Sigmoid

We use the same inputs:

- $x = 2$
- $w = 3$

- $b = 1$

Compute:

$$z = wx + b = (3)(2) + 1 = 7$$

Apply Sigmoid:

$$a = \sigma(7) = \frac{1}{1 + e^{-7}} \approx 0.999$$

Assume the same loss function:

$$L = (a - y)^2, \quad y = 5$$

Step 3: Backpropagation

To update w , compute:

$$\frac{dL}{dw} = \frac{dL}{da} \cdot \frac{da}{dz} \cdot \frac{dz}{dw}$$

1. Compute $\frac{dL}{da}$

$$\frac{dL}{da} = 2(a - y) = 2(0.999 - 5) = -8.002$$

2. Compute $\frac{da}{dz}$ (Sigmoid Derivative)

$$\frac{da}{dz} = \sigma(z)(1 - \sigma(z)) = 0.999(1 - 0.999) = 0.000999$$

3. Compute $\frac{dz}{dw}$

$$\frac{dz}{dw} = x = 2$$

Step 4: Compute Gradient and Update Weights

$$\frac{dL}{dw} = (-8.002) \times (0.000999) \times (2) \approx -0.016$$

Using gradient descent with $\eta = 0.1$:

$$w_{\text{new}} = 3 - 0.1 \times (-0.016) = 3.0016$$

Comparison: ReLU vs. Sigmoid

Feature	ReLU	Sigmoid
Forward Output	$a = 7$	$a \approx 0.999$
Gradient	1	0.000999
Weight Update	$w_{\text{new}} = 2.2$	$w_{\text{new}} = 3.0016$
Gradient Strength	Strong (No Vanishing)	Very weak (Vanishing Gradient)

Key Takeaways

1. **Sigmoid suffers from vanishing gradients** because its derivative is very small when z is large. This results in **slow learning**.
2. **ReLU provides strong gradients** when $z > 0$, making it more effective in deep networks.
3. **Sigmoid squashes values between (0,1)**, which can be useful for probability outputs but not for hidden layers.

Would you like me to extend this example with multiple layers? 🚀

does relu prevent exploding gradient ?

Great question. Let's unpack this clearly and deeply.

🔍 Does ReLU Prevent Exploding Gradient?

✅ Short Answer:

No, ReLU does not prevent exploding gradients by itself.

It helps **avoid vanishing gradients**, but **exploding gradients can still occur**, especially in deep networks or poorly initialized weights.

📘 Let's Understand Why

♦ Exploding Gradient:

Occurs when gradients **accumulate and grow exponentially** during backpropagation, especially in **deep networks**.

Mathematically, during backprop, we multiply many Jacobians (derivatives of each layer):

$$\frac{dL}{dx} = \frac{dL}{da_n} \cdot \frac{da_n}{dz_n} \cdot \frac{dz_n}{da_{n-1}} \dots \frac{dz_2}{da_1} \cdot \frac{da_1}{dz_1} \cdot \frac{dz_1}{dx}$$

If each term > 1 , the overall gradient **blows up**.

What ReLU Does Well

ReLU:

$$f(x) = \max(0, x)$$

$$f'(x) = \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases}$$

✓ Avoids Vanishing Gradients (unlike Sigmoid/Tanh):

- Derivative is 1, not squashed to near-zero
 - Keeps gradients flowing in positive regions
-

⚠ But Why Doesn't It Prevent Exploding Gradients?

Because:

1. **ReLU doesn't cap gradients** — there's no upper bound.
 2. If input weights are too large, **activations and gradients can still explode**.
 3. In **deep ReLU networks**, if weights are poorly initialized or not regularized, the **magnitude of forward activations** can grow layer-by-layer → leading to exploding gradients during backprop.
-

Empirical Example

Let's say:

- You have 100 layers of ReLU.
- You initialize weights $w \sim \mathcal{N}(0, 1)$.
- Input $x = 1$

Forward pass:



$$a = w_1x \rightarrow w_2a \rightarrow w_3a \rightarrow \dots \rightarrow w_{100}a$$

Each multiplication by a weight could double or triple the magnitude → **output explodes** → **so do gradients**.

How to Handle Exploding Gradients (with or without ReLU)

Technique	Description
Weight Initialization	Use He Initialization with ReLU: $\mathcal{N}(0, \sqrt{2/n})$
Gradient Clipping	Clip gradients to a max norm or value during backprop to stop explosions.
Normalization Layers	Use BatchNorm , LayerNorm to stabilize activations & gradients.
Regularization	L2 regularization (weight decay) can help control weight magnitudes.
Architecture Choices	Use residual connections (e.g., in ResNets) to help with gradient flow.



Final Verdict

Question	Answer
Does ReLU prevent vanishing gradients?	 Yes (for $x > 0$)
Does ReLU prevent exploding gradients?	 No (not by itself)

Would you like a code example showing exploding gradients in a deep ReLU net and how He initialization can fix it?

using example and derivation explain why it does not prevent exploding gradient

Absolutely — let's walk through **why ReLU doesn't prevent exploding gradients**, with both:

-  A mathematical derivation
-  A numerical example

Concept Recap: Exploding Gradients

In deep neural networks, during backpropagation, gradients are propagated **back through multiple layers** by **chain rule**:

$$\frac{dL}{dw} = \frac{dL}{da} \cdot \frac{da}{dz} \cdot \frac{dz}{dw}$$

In a deep network:

$$\frac{dL}{dx} = \prod_{l=1}^L \frac{da^{(l)}}{dz^{(l)}} \cdot \frac{dz^{(l)}}{da^{(l-1)}}$$

If each partial derivative is > 1 in magnitude, then their **product explodes**.

Numerical Example with ReLU

Let's build a 3-layer network with ReLU activation and **no biases** for simplicity.

Forward Pass

Let's define:

- Input $x = 1$
- Weights: $w^{(1)} = w^{(2)} = w^{(3)} = 5$
- Each layer:
 $z^{(l)} = w^{(l)} \cdot a^{(l-1)}$
 $a^{(l)} = \text{ReLU}(z^{(l)})$

Layer 1

$$z^{(1)} = 5 \cdot 1 = 5, \quad a^{(1)} = \text{ReLU}(5) = 5$$

Layer 2

$$z^{(2)} = 5 \cdot 5 = 25, \quad a^{(2)} = \text{ReLU}(25) = 25$$

Layer 3

$$z^{(3)} = 5 \cdot 25 = 125, \quad a^{(3)} = \text{ReLU}(125) = 125$$

Loss (mean squared error to target $y = 0$):

$$L = \frac{1}{2}(a^{(3)} - y)^2 = \frac{1}{2}(125)^2 = 7812.5$$

Backpropagation

We want $\frac{dL}{dw^{(1)}}$, so use chain rule:

$$\frac{dL}{dw^{(1)}} = \frac{dL}{da^{(3)}} \cdot \frac{da^{(3)}}{dz^{(3)}} \cdot \frac{dz^{(3)}}{da^{(2)}} \cdot \frac{da^{(2)}}{dz^{(2)}} \cdot \frac{dz^{(2)}}{da^{(1)}} \cdot \frac{da^{(1)}}{dz^{(1)}} \cdot \frac{dz^{(1)}}{dw^{(1)}}$$

Break this down:

1. Output Layer:

$$\frac{dL}{da^{(3)}} = a^{(3)} - y = 125$$

$$\frac{da^{(3)}}{dz^{(3)}} = 1 \quad (\text{ReLU' since } z^{(3)} > 0)$$

$$\frac{dz^{(3)}}{da^{(2)}} = w^{(3)} = 5$$

2. Hidden Layer 2:

$$\frac{da^{(2)}}{dz^{(2)}} = 1 \quad (\text{ReLU'})$$

$$\frac{dz^{(2)}}{da^{(1)}} = w^{(2)} = 5$$

3. Hidden Layer 1:

$$\frac{da^{(1)}}{dz^{(1)}} = 1 \quad (\text{ReLU'})$$

$$\frac{dz^{(1)}}{dw^{(1)}} = x = 1$$

Putting all together:

$$\frac{dL}{dw^{(1)}} = 125 \cdot 1 \cdot 5 \cdot 1 \cdot 5 \cdot 1 \cdot 1 = \boxed{3125}$$

🎯 The gradient for $w^{(1)}$ is **3125**, which is **huge**.



Let's Generalize This

Assume:

- Each weight $w = k$
- Each ReLU passes gradient = 1 (since $z > 0$)
- Network has L layers

Then the gradient becomes:

$$\frac{dL}{dw^{(1)}} = \text{Initial Loss Gradient} \cdot k^{L-1}$$

So for:

- $k = 5, L = 10$: gradient $\propto 5^9 = 1,953,125$ ✅ Explodes

- $k = 1$: gradient stays stable
- $k = 0.5$: gradient vanishes

⚠ Why ReLU Doesn't Fix Exploding Gradient

✗ What ReLU does not do:

- It does **not limit the magnitude** of forward activations or gradients.
- If weight magnitudes are large, forward pass activations **amplify** → so do gradients.

✓ What ReLU does do:

- Keeps gradient alive for $z > 0$ (helps with vanishing gradients).
- Efficient due to sparse activations (0s for negatives).

🔧 Real Fix: Proper Initialization

To avoid exploding gradients in ReLU nets:

Use He Initialization:

$$w \sim \mathcal{N}\left(0, \sqrt{\frac{2}{n_{\text{in}}}}\right)$$

Why? Because it keeps the **variance of activations and gradients** stable across layers.

✓ Summary Table

Concept	ReLU Helps?	Why?
Vanishing Gradient	✓ Yes	Derivative is 1 for $x > 0$
Exploding Gradient	✗ No	No control over magnitude of activations
Needs Proper Initialization	✓ Yes	Use He Initialization
Needs Gradient Clipping	✓ Sometimes	In very deep or recurrent nets

Would you like to visualize exploding vs stable gradients with a deep ReLU net in code (e.g., PyTorch or Keras)?

i did not understand the role of he initialisation

Perfect — let's break down **He initialization** from first principles and show **why** it prevents exploding or vanishing gradients **mathematically and intuitively** using ReLU networks.

What is the Core Problem?

In deep networks, especially with many layers, two things can go wrong during **forward** and **backward** passes:

- ▼ **Vanishing gradients:** gradients shrink → weights stop updating
- ▲ **Exploding gradients:** gradients grow exponentially → training becomes unstable

The root cause: improper scaling of weight initialization → causes activations and gradients to either blow up or vanish across layers.



Let's Derive the Problem

Assume:

- We have a network with **ReLU activations**
- Each layer:

$$z^{(l)} = w^{(l)} \cdot a^{(l-1)} \quad , \quad a^{(l)} = \text{ReLU}(z^{(l)})$$

- The inputs $x \sim \mathcal{N}(0, 1)$
- Weights $w \sim \mathcal{N}(0, \sigma^2)$

? What Happens to the Variance of Activations?

Let's track how the **variance** of activations changes as we move from layer to layer.

Suppose:

- $a^{(l-1)}$ has variance $\text{Var}[a^{(l-1)}] = v$
- We want to compute $\text{Var}[a^{(l)}]$

Since:

$$z^{(l)} = \sum_{i=1}^n w_i a_i^{(l-1)}$$

Then:

$$\text{Var}[z^{(l)}] = n \cdot \text{Var}[w] \cdot \text{Var}[a^{(l-1)}]$$

Now, ReLU activation is applied:

$$a^{(l)} = \text{ReLU}(z^{(l)})$$

And for a Gaussian input $z \sim \mathcal{N}(0, \sigma^2)$, ReLU zeroes out half the values \rightarrow so:

$$\text{Var}[\text{ReLU}(z)] = \frac{1}{2} \text{Var}[z] = \frac{1}{2} n \cdot \text{Var}[w] \cdot \text{Var}[a^{(l-1)}]$$

👉 This recursion happens layer-by-layer!

💣 If Initialization is Wrong

If $\text{Var}[w]$ is too small:

- Variance of activations **shrinks exponentially** \rightarrow ▼ vanishing gradients

If $\text{Var}[w]$ is too large:

- Variance of activations **explodes** \rightarrow ▲ exploding gradients
-

✅ He Initialization to the Rescue

🎯 Goal:

Keep the **variance of activations constant** across layers

\rightarrow i.e. prevent both vanishing and exploding

To keep:

$$\text{Var}[a^{(l)}] \approx \text{Var}[a^{(l-1)}]$$

Set:

$$\frac{1}{2} n \cdot \text{Var}[w] = 1 \Rightarrow \text{Var}[w] = \frac{2}{n}$$

✅ Therefore:

He initialization:

$$w \sim \mathcal{N}\left(0, \frac{2}{n_{\text{in}}}\right)$$

Where:

- n_{in} = number of input units to the layer

This keeps:

- **Forward activations** stable
- **Backward gradients** stable (through chain rule)
- Works **specifically well** for ReLU



Contrast With Other Initializations

Initialization	Formula	Works Well With
Xavier/Glorot	$\mathcal{N}(0, \frac{1}{n_{\text{in}} + n_{\text{out}}})$	Tanh, Sigmoid
He Initialization	$\mathcal{N}(0, \frac{2}{n_{\text{in}}})$	ReLU, LeakyReLU



Simple Code Demo (PyTorch)

```
python

import torch.nn as nn

# He Initialization for ReLU network
def init_weights_he(m):
    if isinstance(m, nn.Linear):
        nn.init.kaiming_normal_(m.weight, nonlinearity='relu')

model = nn.Sequential(
    nn.Linear(256, 128),
    nn.ReLU(),
    nn.Linear(128, 64),
    nn.ReLU(),
    nn.Linear(64, 10)
)

model.apply(init_weights_he)
```



Summary

Concept	Explanation
Why exploding/vanishing happens	Multiplicative gradient chain across layers
Role of He Initialization	Sets weight variance to keep activations' variance constant
Why it's ideal for ReLU	ReLU zeros out half activations → needs compensation via larger init variance (×2)
Variance rule	$\text{Var}[w] = \frac{2}{n_{\text{in}}}$

Would you like a Colab-ready notebook to **visualize how variance explodes or stabilizes across layers** with and without He init?